

GENERAL NOTES

ANDREW S. HARD¹

November 16, 2015

CONTENTS

1	Introduction	3
2	Basic Statistics	3
2.1	Deduction and Induction	3
2.2	Conditional Probability and Bayes' Theorem	3
2.3	Correlation and Covariance	5
3	Data Structures	5
3.1	Lists	5
3.2	Trees	5
3.3	Graphs	5
4	Numerical Methods	5
4.1	Minimization	5
4.2	Regression	5
4.3	Matrix Algebra	5
5	Machine Learning	5
5.1	Introduction	5
5.2	Feature Selection	5
5.3	Regression (Again!)	6
5.4	Clustering	6
5.5	Support Vector Machines	6
5.6	Decision Trees	6
5.7	Lasso Method	6
6	Deep Learning	6
6.1	Overview	6
6.2	Perceptrons	6
6.3	MLP: Multi-Layer Perceptrons	6
6.4	RBM: Restricted Boltzmann Machines	6
6.5	CNN: Convolutional Neural Networks	7
6.6	RNN: Recurrent Neural Networks	7
6.7	LSTM: Long Short Term Memory	7
6.8	Auto-Encoders and Deep Belief Networks	7
6.9	Neural Turing Machines	7
7	Lasso Method	7

LIST OF FIGURES

Figure 1	A Boltzmann Machine diagram	6
----------	-----------------------------	---

LIST OF TABLES

¹ *Department of Physics, University of Wisconsin, Madison, United States of America*

1 INTRODUCTION

As a fifth year graduate student, I have come to the realization that significant portions of my skillset and knowledge base are highly specialized. I intend to pursue a career outside of my field of study (experimental high-energy particle physics). In order to enhance my future job prospects, I decided that it would be useful to review basic concepts in computer science, statistics, mathematics, and machine learning.

2 BASIC STATISTICS

2.1 Deduction and Induction

Deduction: if $A \rightarrow B$ and A is true, then B is true.

Induction: if $A \rightarrow B$ and B is true, then A is more plausible.

2.2 Conditional Probability and Bayes' Theorem

The **conditional probability** of an event A assuming that B has occurred, denoted $P(A|B)$, is given by:

$$P(A|B) = \frac{P(A \cap B)}{P(B)}. \quad (1)$$

Via multiplication:

$$P(A \cap B) = P(A|B)P(B). \quad (2)$$

This makes sense. The probability of A and B is the probability of A *given* B times the probability of B . Or, just as reasonably, the probability of A and B is the probability of B *given* A times the probability of A . As an aside, the equation above can be generalized:

$$P(A \cap B \cap C) = P(A|B \cap C)P(B \cap C) = P(A|B \cap C)P(B|C)P(C). \quad (3)$$

Back to the derivation, since $P(A \cap B) = P(B \cap A)$, we can use Equation 2:

$$P(A|B)P(B) = P(B|A)P(A). \quad (4)$$

And re-arranging gives **Bayes' Theorem**:

$$P(A|B) = \frac{P(A)P(B|A)}{P(B)} = \frac{P(A)}{P(B)}P(B|A). \quad (5)$$

In this equation, $P(A)$ is the **prior probability**, the initial degree of belief in A , or the probability of A before B is observed. It is essentially the probability that hypothesis A is true before evidence is collected. The **posterior probability** $P(A|B)$ is the degree of belief after taking B into account.

Bayes' theorem describes the probability of an event, based on conditions that might be relevant to the event. With the Bayesian probability interpretation the theorem expresses how a subjective degree of belief should rationally change to account for evidence (**Bayesian inference**). Note: in the **Frequentist Interpretation**, probability measures a proportion of outcomes, not probability of belief.

Bayes' theorem can also be interpreted in the following manner: A is some hypothesis, and B is the evidence for that hypothesis.

The equation makes intuitive sense, particularly in the last form in which it is presented. If $P(B) \gg P(A)$, the probability $P(A|B)$ will be small because B *usually* occurs without A occurring. Even if $P(B|A)$ were 100%, the posterior probability would still be small. On the other hand, if $P(A) \gg P(B)$, $P(A|B)$ is high, even if the conditional probability $P(B|A)$ is smaller.

Nate Silver has a discussion of Bayesian statistics in "The Signal And The Noise", p245. Let x be the initial estimate of the hypothesis likelihood (**prior probability**). In the book, x is the initial estimate that your spouse is cheating. y is a conditional probability - the probability of an observation being made given that the hypothesis is true. So y in his case was the probability of underwear appearing conditional on the spouse cheating. z is the probability of an observation being made given that the hypothesis is false. So z was the probability of underwear appearing if the spouse was not cheating. The **posterior probability** is the revised estimate of the hypothesis likelihood. In Silver's case, the likelihood that the spouse is cheating on you, given that underwear was found. The posterior probability is given by the expression:

$$x' = \frac{xy}{xy + z(1 - x)}. \quad (6)$$

In Silver's formulation, the denominator is equivalent to the total probability of an observation being made (total probability of underwear being found). So it is the probability of the hypothesis being true times the associated conditional probability of the observation of underwear plus the probability of the hypothesis being false times the associated conditional probability of the observation.

Note that this method can be applied recursively by substituting x' for x in the expression. We can also come up with a recurrent form of Bayes theorem as formulated in Equation 5 by letting $P(A_{n+1}) = P(A_n|B)$.

The derivation below is from mathworld's article on Bayes' Theorem. In this example, let A and S be sets. Furthermore, let

$$S \equiv \bigcup_{i=1}^N A_i, \quad (7)$$

so that A_i is an event in S, and $A_i \cap A_j = \emptyset \forall i \neq j$. Then:

$$A = A \cap S = A \cap \left(\bigcup_{i=1}^N A_i \right) = \bigcup_{i=1}^N (A \cap A_i) \quad (8)$$

Then we can compare the probabilities using the Law of Total Probability:

$$P(A) = P\left(\bigcup_{i=1}^N (A \cap A_i)\right) = \sum_{i=1}^N P(A \cap A_i) = \sum_{i=1}^N P(A_i)P(A|A_i) \quad (9)$$

Using substitution into Bayes' Theorem (Equation 5), we have a new form:

$$P(A_j|A) = \frac{P(A_j)P(A|A_j)}{\sum_{i=1}^N P(A_i)P(A|A_i)}. \quad (10)$$

What is the meaning of this? The denominator is simply $P(A)$, expressed as the sum of its constituents. So this reduces exactly to Bayes theorem as in Equation 5, with $A \rightarrow A_j$ and $B \rightarrow A$.

2.3 Correlation and Covariance

Correlation and covariance are similar. Both concepts describe the degree to which two random variables or sets of random variables tend to deviate from their expected values in similar ways.

Let X and Y be two random variables, with means μ_X and μ_Y and standard deviations σ_X and σ_Y , respectively. The **covariance** is defined (using $\langle \rangle$ to denote expectation value):

$$\sigma_{XY} = \langle (X - \langle X \rangle)(Y - \langle Y \rangle) \rangle. \quad (11)$$

Similarly, the **correlation** of the two variables is defined:

$$\rho_{XY} = \frac{\langle (X - \langle X \rangle)(Y - \langle Y \rangle) \rangle}{\sigma_X \sigma_Y}. \quad (12)$$

So correlation is dimensionless, while covariance is the multiple of the units of the two variables. **Variance** is simply the covariance of a variable with itself (σ_{XX} is usually denoted σ_X^2 , the square of the standard deviation). The correlation of a variable with itself is always 1, except in the degenerate case where the two variances are zero and the correlation does not exist.

3 DATA STRUCTURES

3.1 Lists

3.2 Trees

3.3 Graphs

4 NUMERICAL METHODS

4.1 Minimization

4.2 Regression

4.3 Matrix Algebra

5 MACHINE LEARNING

5.1 Introduction

This section should contain a general description of data science principles, as well as use cases for different algorithms. Many tools for many problems.

Basic problems in machine learning are classification (labeling) and regression (function estimation).

How to do importance ranking of features? Feature importance, global loss function.

5.2 Feature Selection

Garbage in gives garbage out. Interactions between features. Interactions are not the same as correlations. Filters versus wrappers.




Figure 1: An example of a generic Boltzmann Machine versus a Restricted Boltzmann Machine, from http://www.frontiersin.org/files/Articles/58710/fnins-07-00178-r2/image_m/fnins-07-00178-g001.jpg.

5.3 Regression (Again!)

5.4 Clustering

5.5 Support Vector Machines

5.6 Decision Trees

5.7 Lasso Method

6 DEEP LEARNING

6.1 Overview

Explain the use cases of each type of algorithm.

6.2 Perceptrons

6.3 MLP: Multi-Layer Perceptrons

6.4 RBM: Restricted Boltzmann Machines

<https://www.cs.toronto.edu/~hinton/absps/guideTR.pdf>

"**Restricted Boltzmann machines (RBMs)** have been used as generative models of many different types of data including labeled or unlabeled images (Hinton et al., 2006a), windows of mel-cepstral coefficients that represent speech (Mohamed et al., 2009), bags of words that represent documents (Salakhutdinov and Hinton, 2009), and user ratings of movies (Salakhutdinov et al., 2007). In their conditional form they can be used to model high-dimensional temporal sequences such as video or motion capture data (Taylor et al., 2006) or speech (Mohamed and Hinton, 2010). Their most important use is as learning modules that are composed to form deep belief nets (Hinton et al., 2006a)."

A general Boltzmann machine is a network with symmetric weights fully connecting all input and hidden nodes. The details of the weights and nodes are described later in this section. A restricted Boltzmann machine separates nodes into layers and has no intra-layer connections. General Boltzmann machines, while theoretically providing a more diverse architecture, are difficult to train in practice. The difference between a general Boltzmann Machine and the Restricted Boltzmann Machine is illustrated in Figure 1.

The training procedure for RBMs is more complicated than back-propagation. They typically use a procedure called contrastive divergence (Hinton, 2002). There are several meta-parameters including learning rate, momentum, weight-cost, sparsity target, initial weights, number of hidden units, and mini-batch sizes.

For simplicity, the outline below will consider a two-layer network consisting of the visible unit layer representing the input data and a layer of hidden units called **feature detectors**. Deep neural networks can be trained by successively stacking RBMs, with the hidden layer of the first RBM used as an input for the second. Greedy pairwise training of layers has been demonstrated to be effective in learning deep architectures (reference greedy pair training). Input data are stochastic binary pixels (or continuous variables scaled to $[0, 1]$). Feature detectors are also stochastic binary units. The visible units are connected to the feature detectors using symmetrically weighted connections. A configuration of visible and hidden units has an energy (Hopfield, 1982) given by:

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i a_i v_i - \sum_j b_j h_j - \sum_i \sum_j v_i h_j w_{ij} \quad (13)$$

for $i \in \text{visible}$ and $j \in \text{hidden}$. In this equation, v_i and h_j are the binary states of the visible unit i and hidden unit j , respectively. a_i and b_j are their biases and w_{ij} is the weight between them.

The energy can be used to define a probability of this pairing of visible and hidden units (hence the name "Boltzmann Machine"):

$$p(\mathbf{v}, \mathbf{h}) = \frac{1}{Z} e^{-E(\mathbf{v}, \mathbf{h})} \quad (14)$$

Z is the partition function:

$$Z = \sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} \quad (15)$$

A probability assigned by the network for a visible vector \mathbf{v} can be calculated by summing $p(\mathbf{v}, \mathbf{h})$ over all possible hidden vectors:

$$p(\mathbf{v}) = \frac{1}{Z} \sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})} \quad (16)$$

6.5 CNN: Convolutional Neural Networks

6.6 RNN: Recurrent Neural Networks

6.7 LSTM: Long Short Term Memory

6.8 Auto-Encoders and Deep Belief Networks

6.9 Neural Turing Machines

7 LASSO METHOD